



kubernetes



Agenda

- I. Why the need for Kubernetes
- II. Kubernetes competitors
- III. How to work with Kubernetes
- IV. What is the Kubernetes architecture
- V. How to setup a Kubernetes cluster




κυβερνήτης

Kubernetes



K8s







I. Why the need for **K8s**?



Small Application



Small Application



S




VM

from

What do now?

Small Application



Smart

n



frontend



...

VM


Small Application




/M

What to use for smaller applications?


- Docker compose
- VM
- VPS
- nothing, just bare-metal machine




Enterprise Application




Enterprise Application



Enterprise Application



Enterprise Application



Ent

ion

VM




What do now?



Rec

Se


How can you make sure the application
is up and healthy **all the time**, or with
minimal intervention?




Enterprise Application



How can you make sure all containers
are updated? Do it manually?



Who takes care of



resource management?



security?



scalability?




high availability?



fault tolerance?




load balancing?






or



K8s managed service



Google Kubernetes Engine




Azure Kubernetes Service





Amazon Elastic Kubernetes Service



or



Container orchestration system



HashiCorp

Nomad




Apache
MESOS[™]




RANCHER[®]

BY SUSE




II. K8s competitors?



- Open source, developed by Google released in 2014
- Manage applications in different environments
- Improves reliability
- Reduces the time and resources of DevOps
- Scheduling scaling
- Managing health checks
- declarative model
- Flexibility in multi-cloud and hybrid cloud environments
- Challenges when managing and scaling containers between cloud providers

Docker Swarm




- simplicity with fast deployments
- Docker Engine provides the option of integration with Kubernetes
- You do not need to re-architect your app to adapt to other container orchestrators
- Ideal for smaller deployments
- Automated load balancing within the Docker containers
- No option to connect containers to storage, less user-friendly when it comes to storage-related issues
- Less robust automation capabilities

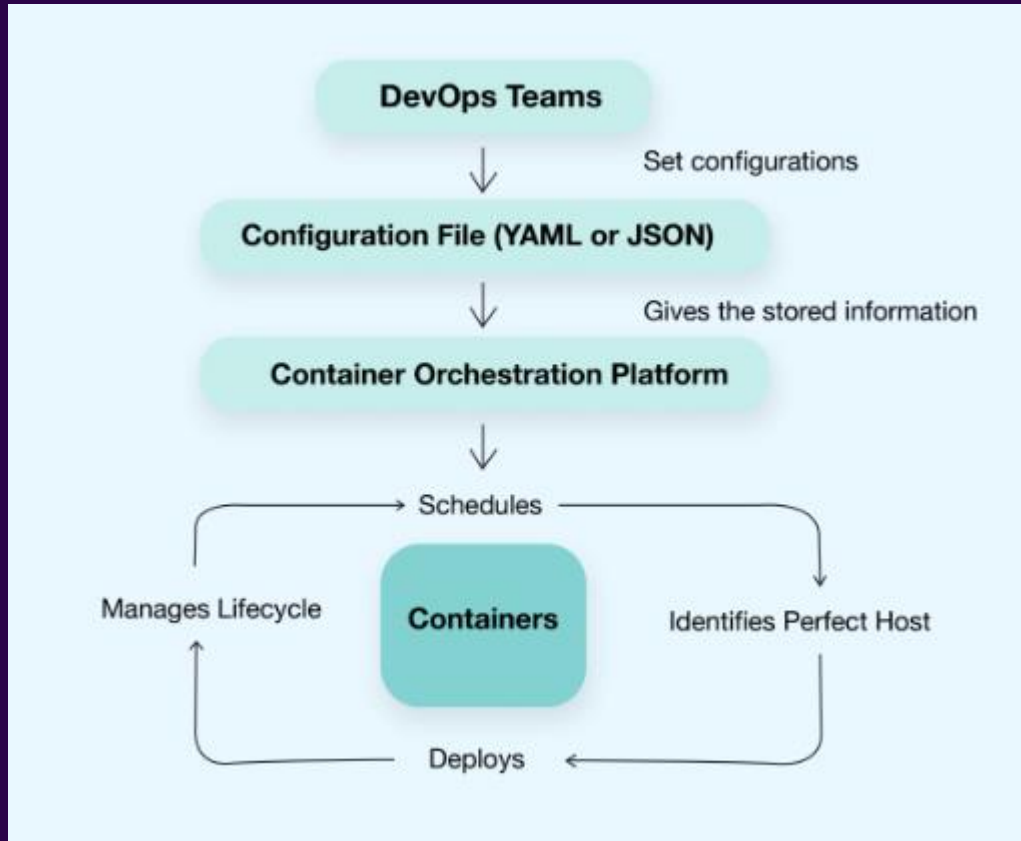
Apache Mesos




- b
- older than Kubernetes
- open-source cluster
- Lightweight interface
- APIs support popular languages like C++, Java, and Python
- only provides management of the cluster -> it is **not a container orchestration system**, therefore, many frameworks have been built on top
- has a high learning curve entails



III. How to **work** with K8s?




1. Requirement for a new container for cluster deployment,
2. K8 schedules an event & identifies the appropriate host based on specifications
3. K8s will manage its lifecycle based on the specifications defined
4. Vertical or horizontal scaling to spread the load across host infrastructure
 - Load balancing
 - Moving containers from one host to another if a host dies



IV. What is the K8s architecture?

Node




A server (physical or virtual)

Node



Nodes



K8s cluster

Master (control plane)


- Controls the cluster

Worker

- Does the work
- Amount depends on work to be done
- Has a „kubelet“ process running on it, makes it possible for the cluster to talk to each other




**What exactly is running on master
node?**

- 
- **API Server** (a container): the only **entrypoint** to cluster
 - **Controller manager**: restarts a container if it dies, keep track of what is happening in the cluster
 - **Scheduler**: scheduling containers on different nodes based on workload
 - **Etcd**: key value storage (holds the current state of the cluster at any time, has all the config data, and all status data of each node and container)
 - **Virtual network**: enables communication between master and worker, turns all the nodes inside of a cluster into one powerful machine that has the sum of all the resources of nodes)




Is one master node enough?

Pod



- Smallest unit in K8s,
- Abstraction over container
- Meant to run 1 application container inside of it




KUBERNETES

DOCKER

**Containerd,
CRI-O,
RunC**

How do they communicate with each other?

- Virtual network
- Each pod gets their own IP (not the container!)
- Pods can die pretty easily ->a new pod gets created and is assigned a new IP address






But if you are communicating with the database pod using the IP address you need to adjust it every time pod restarts...


Service

- A permanent IP that can be attached to a pod
- Lifecycles of service and pod are not connected
- App accesible through web browser? External service




Ingress

- Gives URL a secure protocol and a domain name
- The outside request goes to ingres and the ingres forwards to the service




ConfigMap

- Pods communicate with other other using service
- The communication endpoints are configured inside of the image
- If the endpoint changes you have to rebuild image, push, pull in your pod and restart the whole thing
- **configMap**: contains config data like endpoints and you just connect it to the pod




Secret

- like `configmap` but stores secret data & stored in base64 encoded format
- BUT these secrets are stored **unencrypted** in the API server's **etcd**, anyone with API access and etcd can access them You need 3rd party tools or cloud provider tools to encrypt




Volumes

- K8 does not manage data persistence
- Vol attaches a physical storage on a hard drive to your pod






What happens if my application pod
dies? Crashes or i have to restart it bc
there is a new image



Downtime


Replicas

- Replicate everything on multiple servers (nodes)
- Replica is connected to the same service
- Service is also a load balancer, it will catch the request and forward it to whichever pod is least busy





Deployment

- You will not be creating pods, you will be creating deployments
- Abstraction of pods making it more convenient to replicate them and interact with them





So what would happen now if one of
the replicas dies?







What if the database dies?

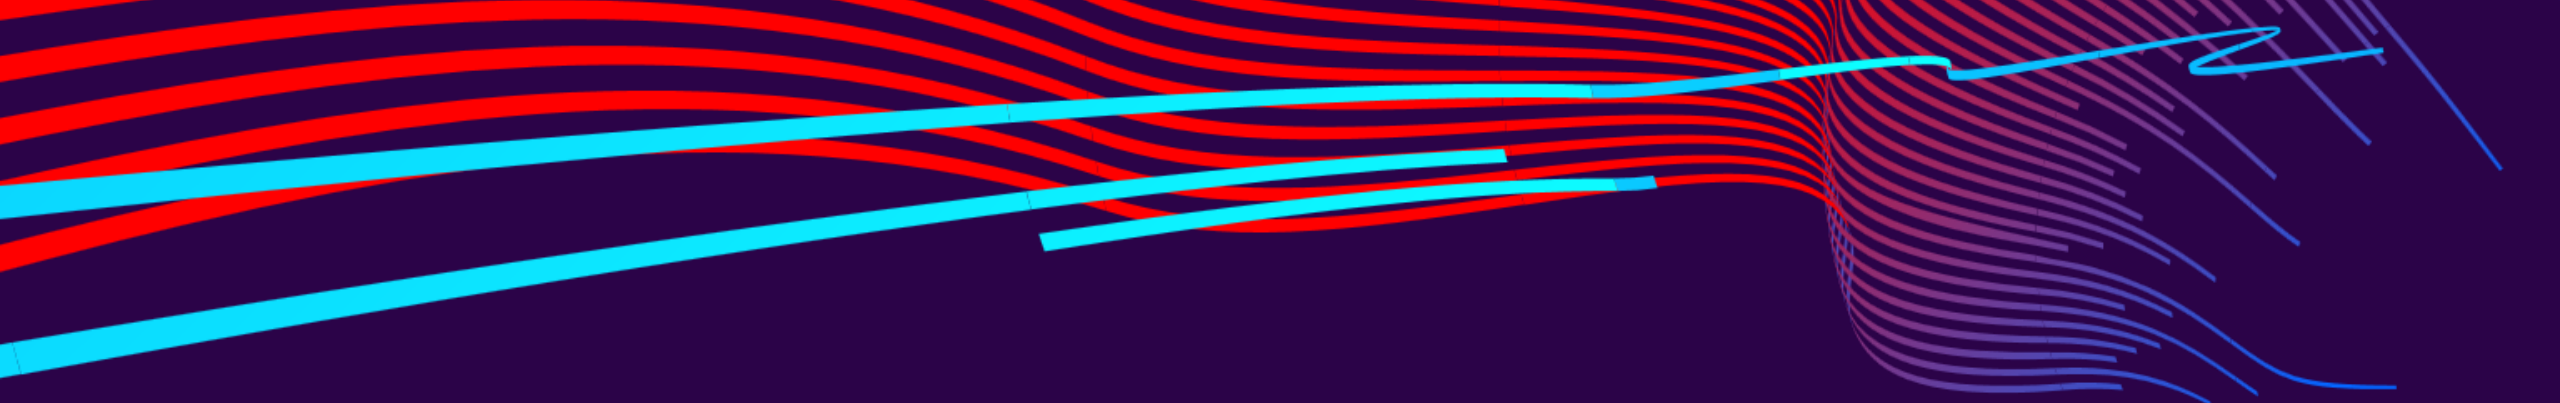
StatefulSet

- The database cannot be replicated using the deployment bc the db has state
- **StatefulSet**: clones access the same shared data storage & avoid data inconsistencies
- Host the database outside K8 cluster and have a stateless app communicating with it





V. How to setup a K8s cluster?

- 
- All the configuration in K8 cluster goes through a master node with the **process API server** where config requests (deployment, service) are sent to.
 - To talk to the API server, you can use
 - UI (K8 dashboard),
 - API (script, curl command),
 - CLI like **kubectl** they all talk to the API server and send their config requests.

K8s configuration file

1) Metadata

where the metadata of that component that you are creating resides

2) Specification

configs you want to apply to that component

3) Status

generated by K8s, K8s will compare what is the desired state in the specification and what is the actual state, K8s updates the state continuously

```
! nginx-deployment.yaml x
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5    labels: ...
7  spec:
8    replicas: 2
9    selector: ...
12   template: ...
22

! nginx-service.yaml x
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: nginx-service
5  spec:
6    selector: ...
8    ports: ...
12
```

Where does k8 get this status data?

- from the `etcd`, the master process that stores the cluster data
- `etcd` holds at any time the current status of any K8s component & this is where the status info comes from

```
status:
  availableReplicas: 1
  conditions:
    - lastTransitionTime: "2020-01-24T10:54:59Z"
      lastUpdateTime: "2020-01-24T10:54:59Z"
      message: Deployment has minimum availability.
      reason: MinimumReplicasAvailable
      status: "True"
      type: Available
    - lastTransitionTime: "2020-01-24T10:54:56Z"
      lastUpdateTime: "2020-01-24T10:54:59Z"
      message: ReplicaSet "nginx-deployment-7d64f4b"
      reason: NewReplicaSetAvailable
      status: "True"
      type: Progressing
  observedGeneration: 1
  readyReplicas: 1
  replicas: 1
  updatedReplicas: 1
```



Who can do the setup?

1. You from hand
2. Tool-supported, e.g., kops (kubernetes operations) with automations
3. Managed cluster – you do no setting up of the cluster, you are using someone elses, e.g., cloud providers
 - You pay for the worker & control plane nodes

**THANK YOU
FOR YOUR
ATTENTION**

REPO:

[GITHUB.COM/ZEZL7/ESD-2024-KUBERNETES](https://github.com/ZEZL7/ESD-2024-KUBERNETES)

